


# Towards Making the Most of ChatGPT for Machine Translation




Keqin Peng<sup>◇,✉\*</sup>, Liang Ding<sup>✉</sup>, Qihuang Zhong<sup>✉</sup>, Li Shen<sup>✉</sup>  
Xuebo Liu<sup>♯</sup>, Min Zhang<sup>♯</sup>, Yuanxin Ouyang<sup>◇</sup>, Dacheng Tao<sup>✉</sup>

<sup>◇</sup>Beihang University <sup>✉</sup>JD Explore Academy <sup>♯</sup>Harbin Institute of Technology, Shenzhen

 keqin.peng@buaa.edu.cn, dingliang1@jd.com

 <https://github.com/Romainpkq/ChatGPT4MT>

## Abstract


ChatGPT shows remarkable capabilities for machine translation (MT). Several prior studies have shown that it achieves comparable results to commercial systems for high-resource languages, but lags behind in complex tasks, e.g., low-resource and distant-language-pairs translation. However, **they usually adopt simple prompts which can not fully elicit the capability of ChatGPT**. In this report, we aim to further mine ChatGPT’s translation ability by revisiting several aspects:  temperature,  task information, and  domain information, and correspondingly propose two (*simple but effective*) prompts: *Task-Specific Prompts* (TSP) and *Domain-Specific Prompts* (DSP). We show that: ❶ The performance of ChatGPT depends largely on temperature, and a lower temperature usually can achieve better performance; ❷ Emphasizing the task information further improves ChatGPT’s performance, particularly in complex MT tasks; ❸ Introducing domain information can elicit ChatGPT’s generalization ability and improve its performance in the specific domain; ❹ ChatGPT tends to generate hallucinations for non-English-centric MT tasks, which can be partially addressed by our proposed prompts but still need to be highlighted for the MT/NLP community. We also explore the effects of advanced in-context learning strategies and find a (*negative but interesting*) observation: the powerful chain-of-thought prompt leads to word-by-word translation behavior, thus bringing significant translation degradation.


## 1 Introduction

Recently, the emergence of ChatGPT<sup>1</sup> has brought remarkable influence on natural language processing (NLP) tasks. ChatGPT is a large-scale language model developed by OpenAI, based on Instruct-GPT (Ouyang et al., 2022), that has been trained

to follow instructions with human feedback. ChatGPT possesses diverse abilities of NLP, including question answering, dialogue generation, code debugging, generation evaluation, and so on (Qin et al., 2023; Zhong et al., 2023; Wang et al., 2023; Kocmi and Federmann, 2023; Lu et al., 2023). We are particularly interested in how well ChatGPT can perform on the machine translation task.

Previous studies (Jiao et al., 2023; Hendy et al., 2023) on translation tasks have found that ChatGPT performs competitively with commercial translation products (e.g., Google Translate and Microsoft Translator) on high-resource languages, but has limited capabilities for low-resource and distant languages. However, they only adopt simple prompts and basic settings regardless of the significant influence of the prompts’ quality (Zhou et al., 2022), which may limit ChatGPT’s performance. In this report, we aim to further elicit the capability of ChatGPT by revisiting the following three aspects and correspondingly propose two simple but effective prompts: *Task-Specific Prompts* (TSP) and *Domain-Specific Prompts* (DSP).


 **Temperature.** Temperature is an important parameter to ensure ChatGPT generates varied responses to human queries. Basically, decoding with higher temperatures displays greater linguistic variety, while the low one generates grammatically correct and deterministic text (Ippolito et al., 2019). However, for tasks with a high degree of certainty, such as machine translation, we argue, a diverse generation may impede its translation quality. We evaluate the performance of ChatGPT at different temperatures to verify its effect and find the optimal temperature setting for the following experiments.

 **Task Information.** ChatGPT is fine-tuned on high-quality chat datasets and thus essentially a conversational system that has a certain distance from the translation system, we argue that the task inconsistency will limit its translation ability to a





\*Work was done when Keqin was interning at JD Explore Academy.

<sup>1</sup><https://chat.openai.com>

certain degree. In response to this problem, we proposed *Task-Specific Prompts* (TSP) to further emphasize the task information to bridge the task gap, i.e., conversation and translation.

 **Domain Information.** Compared with traditional machine translation systems, ChatGPT can incorporate additional information, like human interactions, through the input prompts (Dong et al., 2023). We argue that such flexible interaction may alleviate some classical MT challenges, e.g., cross-domain generalization (Koehn and Knowles, 2017). We, therefore, propose *Domain-Specific Prompts* (DSP) to introduce the domain navigation information to elicit ChatGPT’s generalization ability across different domains.

Through extensive experiments, we find that:

-  ChatGPT’s performance largely depends on the temperatures, especially in difficult languages. Generally, setting a lower temperature can result in higher performance.
-  Emphasizing the task information in prompts can further improve ChatGPT’s performance, especially in complex tasks.
-  Introducing the correct domain information consistently improves ChatGPT’s performance while wrong domain information leads to significant degradation in performance.
-  When tackling the non-English-centric tasks (both the input and expected output are non-English), ChatGPT may generate hallucinations, which should be paid more attention to by the MT/NLP community.

Furthermore, we explore the effects of several advanced in-context learning strategies (Brown et al., 2020). Specifically, we investigate ChatGPT’s few-shot in-context learning (ICL) and chain-of-thought (CoT) (Wei et al., 2022; Kojima et al., 2022) abilities on MT tasks. Experimental results show that few-shot ICL can further improve ChatGPT’s performance, which is identical to the findings of Hendy et al. (2023), and we also find a negative but interesting observation: CoT leads to word-by-word translation behavior, thus bringing significant translation degradation. Also, we call for improving ICL and CoT for MT upon ChatGPT by incorporating the philosophy of example-based and statistical MT (Nagao, 1984; Koehn, 2009).

Test Set	Direction	Domain	Size
Flores-200	Any	General	1,012
WMT19 News	En⇒Zh	News	2,001
	En⇒De		3,004
WMT19 Bio	En⇒Zh	Biomedical	224
	Zh⇒En		241

Table 1: Data statistics and descriptions.

The remainder of this report is designed as follows. We present the evaluation settings in Section 2. In Section 3, we show the zero-shot translation performance of ChatGPT with our proposed advanced prompt recipes. Section 4 summarizes the few-shot in-context learning and chain-of-thought results. Section 5 presents conclusions.

## 2 Evaluation Setting

We provide a brief introduction of the evaluation setting, which mainly includes the used models, test set, and evaluation metrics.

**Models.** We mainly compare ChatGPT<sup>2</sup> with the commercial translation product Google Translator<sup>3</sup>, which supports translation in 133 languages. By default, the results in this report come from the *gpt-3.5-turbo-0301* models, which power the ChatGPT.

**Data.** For multilingual translation and in-context learning, we evaluate the performance of the models on the Flores-200 (Goyal et al., 2022)<sup>4</sup> test sets, which consists of 1012 sentences translated into 204 languages. To evaluate the effect of cross-domain translation, we adopt the test set of WMT19 Biomedical (Bawden et al., 2019) and News Translation Task (Barrault et al., 2019). Table 1 lists the statistics of these test sets. We test all samples through OpenAI API.

**Metric.** The translation metrics shared task (Freitag et al., 2022) recommends using neural network-based metrics since they have demonstrated a high correlation with human evaluation and are resilient to domain shift. Hence, we adopt the mostly used COMET (Rei et al., 2020) as our primary metric. Specifically, we use the reference-based metric COMET-20 (*wmt20-COMET-da*). Additionally, we also report results using SacreBLEU (Post, 2018)

<sup>2</sup><https://chat.openai.com/chat>

<sup>3</sup><https://translate.google.com>

<sup>4</sup><https://github.com/facebookresearch/flores>




Figure 1: The relationship between temperature and ChatGPT’s performance (in terms of COMET scores) when translating from English to other languages.

and ChrF (Popović, 2015) as a reference, but notably, we mainly analyze the performance in terms of model-based metric COMET.

### 3 Zero-Shot Translation

In this section, we explore the performance of ChatGPT from three aspects: TEMPERATURE, TASK INFORMATION, and DOMAIN INFORMATION, and correspondingly propose two simple and effective prompts to improve ChatGPT’s performance.

#### 3.1 The Effect of Temperature 🌡️

ChatGPT is a chatting machine designed to provide fluent and diverse responses to a wide range of human requests. It is intuitive that the diversity of responses may hinder its performance on tasks with a high degree of certainty, such as machine translation, to some extent.

To investigate the influence of diversity, we compare the performance of ChatGPT in different temperature settings, including 0, 0.2, 0.4, 0.6, 0.8, and 1, across three translation directions: English⇒Romanian, English⇒Chinese, and English⇒German. The relationship between temperature and performance of ChatGPT is shown in Figure 1 and 2.

**Results.** Figure 1 and 2 show that ChatGPT’s performance largely depends on the value of temperatures, and as the temperature rises, there is a clear degradation both in COMET and BLEU scores. Furthermore, it is noteworthy that ChatGPT’s sensitivity to the temperature varies depending on the language pair: the impact of temperature




Figure 2: The relationship between temperature and ChatGPT’s performance (in terms of BLEU scores) when translating from English to other languages.

Method	Translation Prompt
ChatGPT	"role": "user", "content": "Please provide the [TGT] translation for the following sentence:"
ChatGPT + TSP	"role": "system", "content": "You are a machine translation system.", "role": "user", "content": "Please provide the [TGT] translation for the following sentence:"

Table 2: Multilingual translation prompts.

is relatively small when translating to high-resource languages, e.g., German, while for complex languages, e.g., Chinese, it has a large degradation in performance (−4.3 COMET points and −3.7 BLEU points for Chinese) when the temperature changes from 0 to 1. We speculate that the huge resource variance in training data leads to differences in the confidence of languages, which partially explains the different performances. In the following experiments, we adopt  $T = 0$  as our default setting to make the most of ChatGPT and ensure the stability of generation.

#### 3.2 The Effect of Task Information 🎯

Previous studies (Jiao et al., 2023; Hendy et al., 2023) have shown that ChatGPT can achieve exceptional performance in conversational domain translation, which is attributed to its ability to generate more natural and diverse spoken language. However, when asking the ChatGPT to perform as a general MT engine, there will arise a task gap. And this task inconsistency may limit ChatGPT’s effectiveness in translation tasks other than the spoken domain.

To bridge the task gap and generate more translation-like sentences, we propose *Task-Specific Prompts* (TSP) to emphasize the transla-

System	COMET	BLEU	ChrF	COMET	BLEU	ChrF
			DE⇒EN	EN⇒DE		
Google Translator	<b>77.7</b>	<b>47.4</b>	<b>70.5</b>	<b>70.5</b>	<b>44.4</b>	<b>68.9</b>
ChatGPT	77.2	43.5	69.4	69.3	<u>40.6</u>	<u>67.1</u>
ChatGPT + TSP	<u>77.5</u>	<u>44.1</u>	<u>69.7</u>	<u>69.4</u>	40.4	67.0
			ZH⇒EN	EN⇒ZH		
Google Translator	<b>73.5</b>	<b>33.5</b>	<b>61.2</b>	<b>68.5</b>	<b>48.8</b>	<b>43.8</b>
ChatGPT	71.3	26.4	58.3	66.4	45.1	39.0
ChatGPT + TSP	<u>71.5</u>	<u>26.7</u>	<u>58.4</u>	<u>67.2</u>	<u>45.3</u>	<u>39.3</u>
			RO⇒EN	EN⇒RO		
Google Translator	<b>82.4</b>	<b>48.0</b>	<b>71.2</b>	91.6	<b>43.3</b>	<b>67.0</b>
ChatGPT	80.6	41.8	68.8	92.4	40.6	65.5
ChatGPT + TSP	<u>80.8</u>	<u>41.9</u>	<u>69.0</u>	<b>92.9</b>	<u>40.8</u>	<u>65.7</u>
			ZH⇒RO	RO⇒ZH		
Google Translator	73.9	<b>25.8</b>	<b>53.9</b>	<b>62.3</b>	<b>42.3</b>	<b>37.8</b>
ChatGPT	73.8	20.9	<u>51.5</u>	58.9	37.7	33.3
ChatGPT + TSP	<b>74.1</b>	<u>21.0</u>	51.3	<u>59.1</u>	<u>38.0</u>	<u>33.7</u>

Table 3: Performance with different prompts on 8 language pairs from Flores-200. “TSP” denotes our proposed task-specific prompting method. The best scores across different systems are marked in **bold** and the best scores of ChatGPT are underlined. Notably, we set the temperature as 0 for ChatGPT in this experiment. We can see that our TSP method consistently boosts the performance of ChatGPT in most settings. **Shadowed** areas mean difficult English-centric translation tasks, **Green** areas mean non English-centric translation tasks.




Figure 3: Number of Post-Edited sentences in non-English-centric language pairs, where a higher value means the translation contains more hallucinations. RO represents the translation for ZH⇒RO, while ZH represents the translation for ZH⇒RO.

tion task information. Specifically, we prepend the sentence “You are a machine translation system.” to the best translation template in Jiao et al. (2023), and adopt it to query ChatGPT. The templates of prompts present in Table 2, and [TGT] represents the target languages of translation.

We have compared the performance of various models on four language pairs, covering eight dis-

tinct translation directions. These languages comprise 1) German, which is one of the most non-English languages in the GPT training data, 2) Romanian, a less frequently encountered non-English language in the GPT training data, and 3) Chinese, a large-scale language with a script distinct from English. We also adopt Chinese-Romanian as a non-English-centric use case. Table 3 lists the full results, where we list both English-centric and non-English-centric language directions (marked with **green**), and also, among English-centric directions, we highlight the difficult pairs (EN-ZH and EN-RO with **shadow**) in terms of their resources and language distance.


### 3.2.1 English-Centric Language Pairs

We first consider the performance of ChatGPT in English-centric translation language pairs. Specifically, we conduct experiments in three language pairs: German⇒English (high-resource), Romanian⇒English (low-resource), and Chinese⇒English (distant language).



**Results.** Our results presented in Table 3 show that our TSP method achieves comparable results on COMET score compared to Google Translator and even outperforms it in some language pairs, e.g., English⇒Romanian (92.9 v.s. 91.6). We also observe that our TSP method consistently improves the performance of vanilla ChatGPT, especially when translating to low-resource or distant languages. Specifically, our TSP method brings +0.8 and +0.5 COMET score improvements in English⇒Chinese and English⇒Romanian, respectively, and +0.2 on average when translating to English. We speculate that the high-resource training data can help the model better understand the specific task from a few task-related navigations, thereby reducing the need for additional task-specific information. Although our proposed TSP consistently improves the performance in terms of semantic metric, i.e., COMTE, notably, we have not consistently bridged the task gap in terms of lexical metrics (BLEU and ChrF).

### 3.2.2 Non-English-Centric Language Pairs

We also evaluate the performance of ChatGPT in non-English-centric language pairs. We have an important finding that,  **when tackling non-English-centric MT language pairs, ChatGPT tends to generate hallucinations**, that is, some unrelated information obeyed some patterns followed the translation, such as *"Translation may vary depending on context"*, which will greatly affect the MT performance. We used a post-processing method to remove irrelevant information from the generated text. The number of post-processed sentences is presented in Figure 3.

**Results.** Figure 3 shows that lower temperature can reduce the number of hallucinations, and our TSP method can further reduce its number, which suggests that our method can help ChatGPT to better serve as a machine translation system. The full results on Romanian⇌Chinese lists are in Table 3. As seen, our TSP method can only slightly improve ChatGPT’s performance, which could be due to the difficulty in both understanding and generating the language pairs. Although our used post-editing approach could remove the hallucination patterns, *the NLP/MT community should pay more attention to the potential hallucination when using ChatGPT to tackle the non-English text.*

The subsequent experiments will use ChatGPT with TSP as the default setting.

Method	Translation Prompt
ChatGPT	"role": "system", "content": "You are a machine translation system.", "role": "user", "content": "Please provide the [TGT] translation for the following sentence: "
ChatGPT+DSP	"role": "system", "content": "You are a machine translation system that translates sentences in the [DOM] domain.", "role": "user", "content": "Please provide the [TGT] translation for the following sentence: "
ChatGPT+F-DSP	"role": "system", "content": "You are a machine translation system that translates sentences in the [FDOM] domain.", "role": "user", "content": "Please provide the [TGT] translation for the following sentence: "

Table 4: Domain-Specific translation prompts. “[DOM]” and “[FDOM]” denote the correct and incorrect domain instructions, respectively.

### 3.3 The Effect of Domain Information

Compared with traditional machine translation systems, ChatGPT can incorporate additional information through the prompts to further improve its performance. While previous studies have shown that ChatGPT has great robust translation capabilities (Hendy et al., 2023), we believe that we can further enhance its performance by incorporating domain-specific guidance.

To this end, we propose Domain-Specific Prompts (DSP) that identify the domain information of translated sentences in prompts to facilitate ChatGPT’s generalization. Specifically, we ask ChatGPT with the following prompts, as shown in Table 4. Here, [DOM] represents the correct domain of the translated sentence (e.g. news, biomedical), while [FDOM] represents the wrong domain of that.

We evaluate the WMT19 Bio and News datasets, which allow us to examine the impact of domain bias. For example, the WMT19 Bio test set comprises Medline abstracts that require domain-specific knowledge, while the WMT19 News dataset features news-style texts that are significantly different from dialogues.

**Results.** The results are listed in Table 5. Obviously, the original ChatGPT does not perform as well as Google Translator in both COMET and lexical metrics (e.g., BLEU). However, our DSP method can consistently improve the performance of ChatGPT in terms of COMET score, and even outperforms Google Translator in two datasets (WMT19 Bio Chinese ⇒ English and WMT19

System	WMT19 Bio				WMT19 News			
	EN $\Rightarrow$ ZH		ZH $\Rightarrow$ EN		EN $\Rightarrow$ ZH		EN $\Rightarrow$ DE	
	COMET	BLEU	COMET	BLEU	COMET	BLEU	COMET	BLEU
Google Translator	<b>59.4</b>	<b>38.8</b>	59.4	<b>36.1</b>	59.3	<b>43.4</b>	<b>64.1</b>	<b>33.7</b>
ChatGPT	58.6	35.5	58.7	31.1	58.8	39.6	63.1	31.3
ChatGPT + DSP	<u>58.9</u>	<u>35.8</u>	<u>59.6</u>	<u>31.3</u>	<u>59.6</u>	<u>39.8</u>	<u>63.2</u>	<u>31.5</u>
ChatGPT + F-DSP	58.6	35.6	<b>58.4</b>	<u>31.3</u>	<b>57.9</b>	39.0	<b>62.0</b>	31.2

Table 5: Performance of ChatGPT on translation robustness, i.e., different domains. “DSP” denotes our proposed domain-specific prompting method, while “F-DSP” denotes the false domain-specific prompting, i.e., we specify wrong/unrelated domain information in the prompt. The results in **green** denote that “DSP” improves ChatGPT by a clear margin (0.5 ( $\uparrow$ ) score), while the **red** results denote the significant performance drops caused by “F-DSP”.

System	EN $\Rightarrow$ DE		EN $\Rightarrow$ ZH		EN $\Rightarrow$ RO	
	COMET	BLEU	COMET	BLEU	COMET	BLEU
Google Translator	<b>70.5</b>	<b>44.4</b>	<b>68.5</b>	<b>48.8</b>	91.6	<b>43.3</b>
ChatGPT	69.4	40.4	67.2	45.3	92.9	40.8
<i>Standard few-shot prompting</i>						
-w/ 1-shot	69.9	40.5	<u>68.2</u>	45.6	93.4	<u>41.2</u>
-w/ 3-shot	69.7	<u>40.6</u>	68.1	<u>45.7</u>	<b>93.5</b>	41.1

Table 6: Few-shot translation performance of ChatGPT on 20 samples randomly selected from Flores-200. In the standard few-shot prompting setting, we randomly sample 1/3 examples from the development set.

News English  $\Rightarrow$  Chinese). This finding indicates that our method can further improve the generalization ability of ChatGPT and narrow the gap with one of the most advanced commercial systems – Google Translator. Nonetheless, our method’s impact on BLEU is inconsistent, and it still lags significantly behind Google Translator’s performance.

To verify that the observed improvement is indeed due to the introduction of the domain information, we deliberately provided incorrect domain information for each sentence, namely F-DSP, to attack the improvement brought by the DSP strategy. Specifically, We exchange domain information for the biomedical sentences and the news sentences. We expect that the wrong domain guidance (F-DSP) will under-perform the DSP, and even perform worse than the vanilla ChatGPT. The results of these experiments are shown in the last row of Table 5, which clearly shows a consistent degradation in COMET.

All the above DSP and F-DSP results confirm the importance of domain-specific prompting guidance in using ChatGPT for MT tasks.

## 4 Few-shot Machine Translation

In this section, we simply explore the effects of advanced in-context learning (ICL) strategies, specifically, we investigate ChatGPT’s few-shot ICL and Chain-of-Thought (CoT) abilities on MT tasks.

### 4.1 Few-Shot In-Context learning

In-context learning (Brown et al., 2020) has shown its remarkable ability for many NLP tasks (Liu et al., 2023). To further explore the capabilities of the ChatGPT, we conduct experiments with different sample selection strategies. Specifically, we evaluate the performance of few-shot machine translation in the following three directions: English $\Rightarrow$ Chinese, English $\Rightarrow$ Romanian, and English $\Rightarrow$ German in Flores-200. We conducted experiments mainly with randomly sampled demonstrations from development sets in the 1-shot and 3-shot settings.

**Results.** Our results are listed in Table 6. As seen, in-context learning with random examples consistently improves the performance in both lexical metric (BLEU) and COMET score compared to the zero-shot approach, which is consistent with previous finding (Hendy et al., 2023). Furthermore,

Method	Translation Prompt
<b>Zero-Shot CoT</b>	"role": "system", "content": "You are a machine translation system.", "role": "user", "content": "Please provide the German translation for the following sentence step by step and then provide the complete sentence: '
<b>1-Shot CoT</b>	"role": "system", "content": "You are a machine translation system.", "role": "user", "content": "Please provide the German translation for the following sentence step by step and then provide the complete sentence: [S] 1. [S_1] - [T_1] 2. [S_2] - [T_2] ... n. [S_n] - [T_n] The complete sentence in [TGT] is: [T] Please provide the German translation for the following sentence step by step and then provide the complete sentence: '

Table 7: The templates of Zero-Shot CoT and 1-shot CoT. [S<sub>n</sub>] represents the *n*-th token in source demonstration [S], [T<sub>n</sub>] represents the *n*-th token in target demonstration [T].

we observed that the 1-shot approach achieved decent results, but further increasing the number of shots does not lead to any improvement.

We encouragingly find that the advanced sample-selection strategies for in-context learning for MT tasks upon ChatGPT are extremely similar to the design philosophy of example-based machine translation (EBMT, Nagao, 1984), where the EBMT is often characterized by its use of a bilingual corpus as its main knowledge base, at run-time. It is worthy of designing better ICL strategies inspired by EBMT in future work.

## 4.2 Chain-of-Thought

Chain-of-Thought (CoT) prompting (Wei et al., 2022) has been demonstrated to be effective in eliciting the reasoning ability of large language models. Previous studies have shown that CoT can improve the ChatGPT’s performance in natural language understanding tasks (Zhong et al., 2023), but **its influence on machine translation tasks has hardly been investigated.**

To investigate this further, we randomly select 20 samples from the test set and adopt the zero-shot CoT technique (Kojima et al., 2022) and the 1-shot CoT technique. Specifically, as shown in Table 7, for zero-shot CoT, we use the prompt "Please provide the [TGT] translation for the following sentence step by step" to extract step-by-step translation. We also add the sentence

Method	EN⇒DE		EN⇒ZH	
	COMET	BLEU	COMET	BLEU
ChatGPT	72.4	36.5	68.3	41.4
-w zero-shot CoT	69.3 (↓3.1)	35.1 (↓1.4)	59.5 (↓8.8)	36.2 (↓5.2)
-w 1-shot CoT	69.6 (↓2.8)	37.0 (↑0.5)	61.1 (↓7.2)	37.6 (↓3.8)

Table 8: Performance of ChatGPT equipped with CoT prompting methods on English⇒German and English⇒Chinese.

‘and then provide the complete sentence:’ to the end of the prompting to ensure that ChatGPT can generate the complete translation. While for the 1-shot CoT, we provide the manual intermediate reasoning steps inspired by zero-shot CoT, as shown in Table 7.

**Results.** We conduct experiments in the following two translation directions: English⇒German and English⇒Chinese. The results are listed in Table 8, which shows that there is a significant degradation in COMET score with zero-shot CoT setting, especially in English⇒Chinese, which drops 8.8 COMET points. 1-shot CoT prompting can consistently outperform zero-shot CoT but still lags behind zero-shot prompting on COMET.

We looked in detail at the sentences generated by different prompts, presented in Table 9, and we have a negative but interesting observation: the CoT prompt leads to word-by-word translation behavior, which is the main reason for the significant translation degradation.

For more CoT variants designed with different principles inspired by the philosophy in statistical MT (Zens et al., 2002; Koehn, 2009) will be explored in the future. For example, word-by-word and then reordering the translation, phrase-to-phrase and then reordering the translation, and structure-to-structure translation.

## 5 Conclusion

We investigate how to further mine ChatGPT’s translation ability from three perspectives, namely temperature, task & domain information, and we correspondingly propose two simple but effective prompts. Through quantitative studies, we summarize some interesting findings and useful recipes to make the most of ChatGPT for MT.

In future work, besides the aforementioned explorations (EBMT-inspired prompts designing, statistical MT-inspired chain-of-thought designing), we would like to investigate the ability of ChatGPT for more MT settings, e.g., document translation.

## Limitations

Our work has several potential limitations. First, we only propose some simple prompts that have not been carefully designed to investigate the capabilities of ChatGPT, which may not sufficiently elicit the power of ChatGPT. Second, we have not fully studied the performance of ChatGPT in few-shot scenarios, especially the effect of Chain-Of-Thought in machine translation. In future work, we would like to design different types of prompts to further improve ChatGPT's performance in machine translation and conduct more in-depth analyses and discussions.

## References

- Loïc Barrault, Ondrej Bojar, Marta R. Costa-jussà, et al. 2019. [Findings of the 2019 conference on machine translation \(WMT19\)](#). In *WMT*.
- Rachel Bawden, Kevin Bretonnel Cohen, Cristian Grozea, et al. 2019. [Findings of the WMT 2019 biomedical translation shared task: Evaluation for MEDLINE abstracts and biomedical terminologies](#). In *WMT*.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, et al. 2020. [Language models are few-shot learners](#). In *NeurIPS*.
- Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Zhiyong Wu, Baobao Chang, Xu Sun, Jingjing Xu, Lei Li, and Zhifang Sui. 2023. [A survey on in-context learning](#). *arXiv preprint*.
- Markus Freitag, Ricardo Rei, Nitika Mathur, et al. 2022. [Results of WMT22 metrics shared task: Stop using BLEU – neural metrics are better and more robust](#). In *WMT*.
- Naman Goyal, Cynthia Gao, Vishrav Chaudhary, Peng-Jen Chen, et al. 2022. [The flores-101 evaluation benchmark for low-resource and multilingual machine translation](#). *TACL*.
- Amr Hendy, Mohamed Abdelrehim, Amr Sharaf, et al. 2023. [How good are gpt models at machine translation? a comprehensive evaluation](#). *arXiv preprint*.
- Daphne Ippolito, Reno Kriz, João Sedoc, Maria Kustikova, and Chris Callison-Burch. 2019. [Comparison of diverse decoding methods from conditional language models](#). In *ACL*.
- Wenxiang Jiao, Wenxuan Wang, Jen-tse Huang, Xing Wang, and Zhaopeng Tu. 2023. [Is chatgpt a good translator? a preliminary study](#). *arXiv preprint*.
- Tom Kocmi and Christian Federmann. 2023. [Large language models are state-of-the-art evaluators of translation quality](#). *arXiv preprint*.
- Philipp Koehn. 2009. *Statistical machine translation*. Cambridge University Press.
- Philipp Koehn and Rebecca Knowles. 2017. [Six challenges for neural machine translation](#). In *WMT*.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. [Large language models are zero-shot reasoners](#). In *NeurIPS*.
- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2023. [Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing](#). *ACM Comput. Surv.*
- Qingyu Lu, Baopu Qiu, Liang Ding, Liping Xie, and Dacheng Tao. 2023. [Error analysis prompting enables human-like translation evaluation in large language models: A case study on chatgpt](#). *arXiv preprint*.
- Makoto Nagao. 1984. [A framework of a mechanical translation between japanese and english by analogy principle](#). *Artificial and human intelligence*.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, et al. 2022. [Training language models to follow instructions with human feedback](#). *arXiv preprint*.
- Maja Popović. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). In *WMT*.
- Matt Post. 2018. [A call for clarity in reporting BLEU scores](#). In *WMT*.
- Chengwei Qin, Aston Zhang, Zhuosheng Zhang, Jiaao Chen, Michihiro Yasunaga, and Diyi Yang. 2023. [Is chatgpt a general-purpose natural language processing task solver?](#) *arXiv preprint*.
- Ricardo Rei, Craig Stewart, Ana C Farinha, and Alon Lavie. 2020. [COMET: A neural framework for MT evaluation](#). In *EMNLP*.
- Jiaan Wang, Yunlong Liang, Fandong Meng, Zhixu Li, Jianfeng Qu, and Jie Zhou. 2023. [Cross-lingual summarization via chatgpt](#). *arXiv preprint*.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed H. Chi, Quoc Le, and Denny Zhou. 2022. [Chain of thought prompting elicits reasoning in large language models](#). *arXiv preprint*.
- Richard Zens, Franz Josef Och, and Hermann Ney. 2002. [Phrase-based statistical machine translation](#). In *KI*.
- Qihuang Zhong, Liang Ding, Juhua Liu, Bo Du, and Dacheng Tao. 2023. [Can chatgpt understand too? a comparative study on chatgpt and fine-tuned bert](#). *arXiv preprint*.
- Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. 2022. [Large language models are human-level prompt engineers](#). *arXiv preprint*.




---

### Zero-shot prompting

**System:** You are a machine translation system.

**User:** Please provide the German translation for the following sentence: Sadly, as newer methods of writing have emerged, the pencil has been relegated to lesser status and uses.


 **Response:** Leider wurde der Bleistift durch neuere Schreibmethoden in eine geringere Position und Verwendung gedrängt.

---

### Zero-shot CoT prompting

**System:** You are a machine translation system.

**User:** Please provide the German translation for the following sentence step by step and then provide the complete sentence: Sadly, as newer methods of writing have emerged, the pencil has been relegated to lesser status and uses.


 **Response:** 1. Sadly - Bedauerlicherweise 2. as newer methods of writing have emerged - da neuere Schreibmethoden aufgetaucht sind 3. the pencil - der Bleistift 4. has been relegated - wurde zurückgestuft 5. to lesser status and uses - zu geringerem Status und Verwendungszwecken Bedauerlicherweise wurde der Bleistift aufgrund neuerer Schreibmethoden zu geringerem Status und Verwendungszwecken zurückgestuft.

---

### 1-shot CoT prompting

**System:** You are a machine translation system.

**User:** Please provide the German translation for the following sentence step by step and then provide the complete sentence: That said, expect to be out of breath, and take care in the steeper portions, especially when wet, as it can become dangerous quickly. 1. That said - jedoch damit 2. expect - Rechnen sie 3. to be - zu sein 4. out of breath - außer Atem 5. and - und 6. take care - vorsichtig 7. in the steeper portions - in den steileren Abschnitten 8. especially - besonders 9. when wet - wenn es nass ist 10. become - werden 11. dangerous - gefährlich 12. quickly - schnell The complete sentence in German is: Rechnen Sie jedoch damit, außer Atem zu sein, und seien Sie in den steileren Abschnitten vorsichtig, besonders wenn es nass ist, da es dort schnell gefährlich werden kann. **Please provide the German translation for the following sentence step by step and then provide the complete sentence:** Sadly, as newer methods of writing have emerged, the pencil has been relegated to lesser status and uses.

 **Response:** 1. Sadly - Leider 2. as - da 3. newer methods - neuere Methoden 4. of writing - des Schreibens 5. have emerged - aufgetaucht sind 6. the pencil - der Bleistift 7. has been relegated - wurde zurückgestuft 8. to lesser status - zu einem geringeren Status 9. and uses - und Verwendungen. The complete sentence in German is: Leider wurde der Bleistift aufgrund des Auftauchens neuerer Methoden des Schreibens zu einem geringeren Status und Verwendungen zurückgestuft.

---

Table 9: Examples of zero-shot and 1-shot CoT prompting on English⇒German. The sentences in **red** are the reasoning step for CoT and those in **blue** are the one-shot example. Sentences in **bold** are the instruction of CoT.